

## → Canonical Interview

### linux Operating System

#### 1) KERNELS

- \* Core component of linux OS.
- \* Manages hardware and software interactions.
- \* enables multitasking and resource management.

#### kernel subsystem

- Process Management ✓
- Memory Management ✓ (MMU)
- Virtual file system (VFS)
- Network Stack. ✓
- Device drivers. ✓
- IPC unit.

#### Q. What is a kernel Module?

Ans: A piece of code that can be loaded and unloaded into kernel upon demand.

extend functionality without rebooting systems

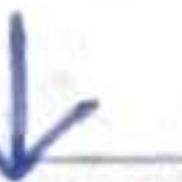
eg: device drivers.

→ lsmod

list the module in kernel.

~~Imp~~ Process of loading.

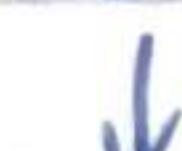
(a) not present feature



(b) Search for it using name.  
(modprobe)



(c) Check if it needs dependencies.



(d) load dependency modules



(e) then load main module.

\* Ubuntu Core

### Kernel Arch

Kernel

↳ Kernel Arch

→ Central Program managing all hardware and software resources.

→ Provides services to user program.

eg: file, scheduling, memory management

### Monolithic kernel (Linux)

→ Most services run in kernel space rather than as a separate processes.

→ Linux also supports loadable kernel module  
↳ add/remove functionality (like drivers)

Show of kernel is structured & different subsystems interconnect.

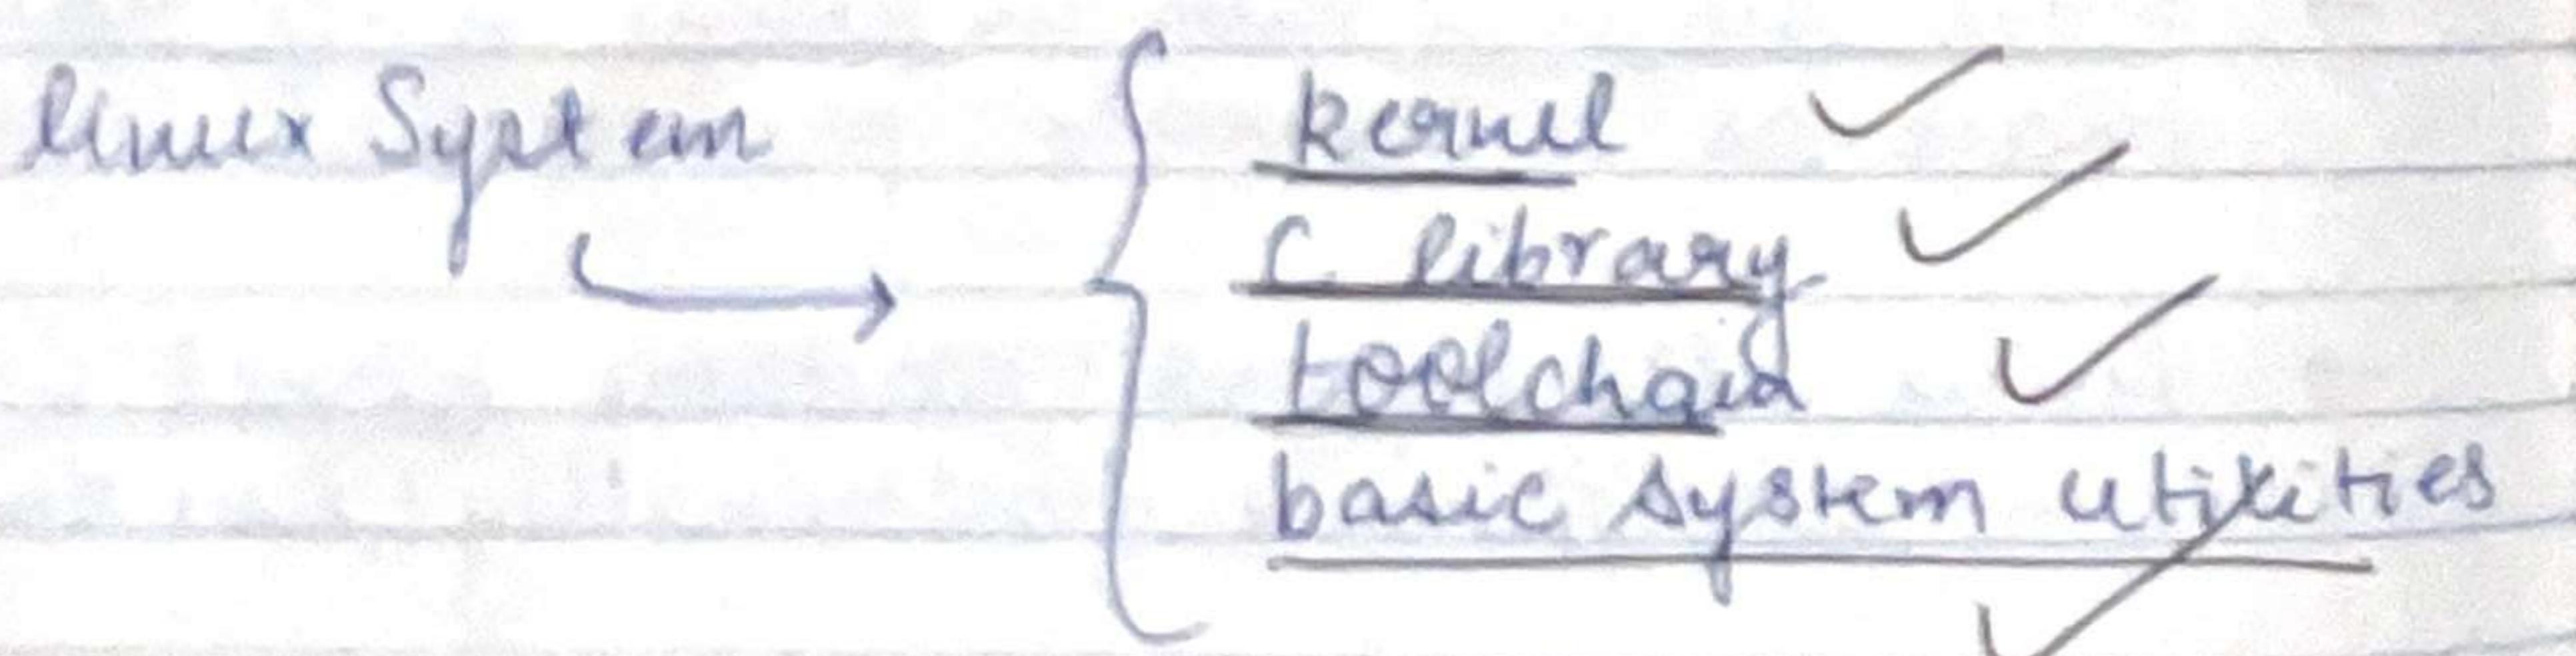
the lower renal artery will

# Introduction to Linux Kernel

- \* Everything is a file. It works in C making it most highly portable.
  - \* unique fork() system call
  - \* has only 100's of system calls.

→ Linux (1991)

Linus Torvalds for Intel 80386



Linux → Refers to → Kernel.

Kernel : Inner most layer of OS

→ Components of kernel

- 1) Interrupt Handler ISR
  - 2) Service Interruption Queue
  - 3) Scheduler
  - 4) Memory management system
  - 5) System Services
    - ↳ Networking
    - ↳ IPC

1wb

Kernel manages system Resources.

## Hardware

↑ Interrupt id.

Intergenetic  
transduc<sup>n</sup>  
processes  
kernel tells no

} Interrupts run in interrupt context  
and not process or context

This exists only to quickly handle  
interrupts

## CPU

A processor at any point is doing 1 of 3 things.

- \* In user-space, executing user code.
- \* In kernel-space, process context interrupt.
- \* In kernel-space, context

eg: idle process

↳ kernel space, process context.

~~Linux~~

Linux usually requires MMU, but it has the capability to run MMU-less embedded systems.

Nowadays (ES) also have MMU.

Complex

Linux is monolithic

- executes in single address space
- entirely in kernel mode.
- Kernel preemption
- Kernel threads.
- dynamically load separate binaries (kernel modules)
- kernel is schedulable.

~~Linux~~

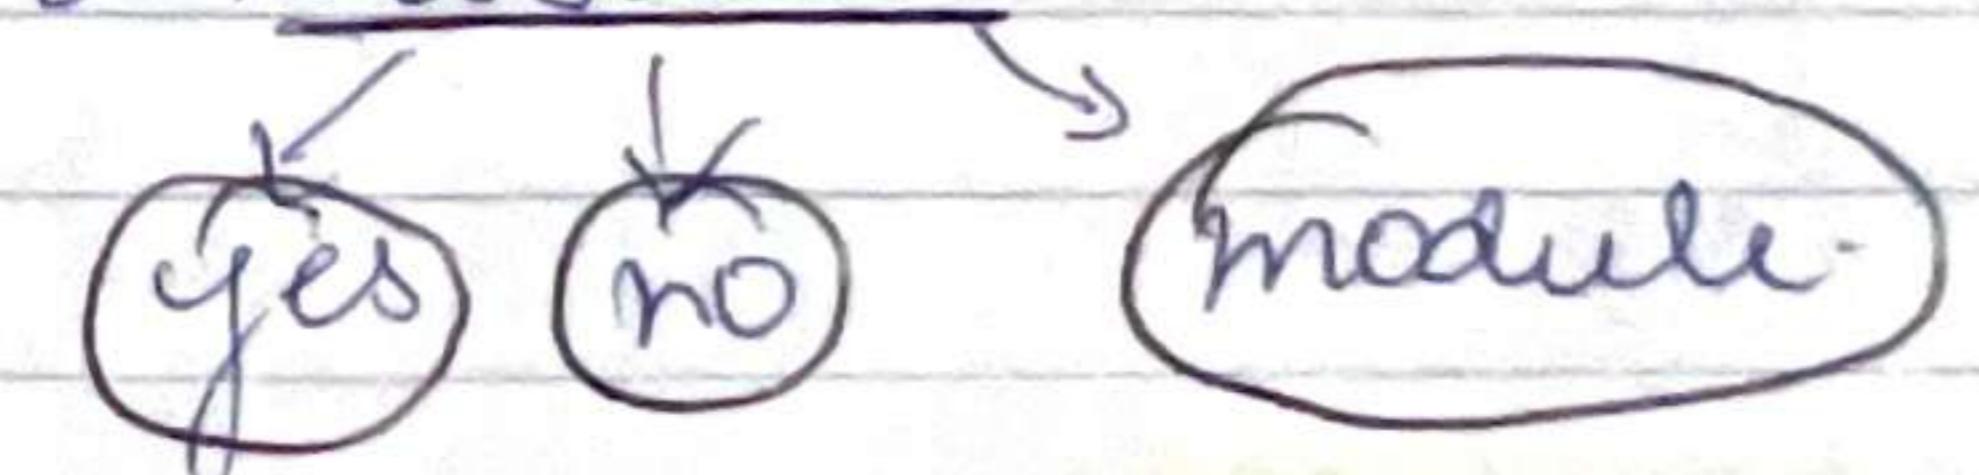
The kernel source is usually installed in

usr/src/linux

usr/src/linux

Should remain untouched

→ configuration option are either boolean or tristate.



Drivers are usually tristate.

Ubuntu kernel

Cross section of kernel features and compile all the drivers as modules

/proc/config.gz

Compressed configuration file is stored here.

build it with Make.

Make file handles dependencies.

→ To build kernel with multiple make jobs.

make -jn.

/boot : store the kernel

/lib/modules : have all installed modules.

→ Extra characteristics.

- 1) Kernel is coded in C.
- 2) Can not easily execute floating up.
- 3) Small per process fixed size stack.

linux → parallelism  
→ concurrency.

\* Kernel memory is not pageable

\* Memory violation in kernel can result in major kernel error.

→ linux supports symmetrical multiprocessing

↳ preemptive multitasking operating System.

### Process Management

→ linux uses a "completely fair scheduler" (CFS)

↳ based on (a) Priority  
(b) time slicing RR

### Real-time Scheduling

↳ FIFO      ↳ RR.

### Memory Management

linux uses virtual memory

- paging
- swapping
- allocation

~~ISR~~

~~IDT~~

## Interrupts

→ Kernel has (ISR) to handle interrupts

↳ interrupts are high priority.



Interrupt Vector.  
(table to identify the interrupt)

Top half & bottom half

↳ minimal work done upon interrupt

Work done later ←  
that requires extensive processing.

## Kernel Boot Process

1) Firmware / BIOS / UEFI

↳ initialises the hardware and loads the bootloader

2) Bootloader (GRUB, LILO)

↳ loads kernel image & an optional initial ram disk

## Kernel Initialisation

↳ kernel decompresses

↳ sets up basic hardware and memory.

↳ Starts first user space process

(init or Systemd)

handles services and user-level start up scripts.

until system is fully operational



## Process Management

Struct task\_struct (PCB)

↳ holds process meta data.  
(PID, state, priority etc.)

### Process States

Task-running → READY

Task-Interruptable → WAITING-STATE

CREATION

fork()    exec()

↳ fork() → duplicate of parent process.

[clone() sys call]

copy-on-write

exec() → replaces the process image  
with new executable.

### Process Scheduling

Completely Fair Scheduler (CFS)

2.6.23

↳ uses red-black tree to track  
runnable task.

red-black tree

→ Allocate CPU time & to task priority.

### System Calls

→ user invokes software interrupt.

→ system call number passed via register

### Kernel Data Structure

stored in /proc

linked list

circular, doubly    struct list\_head

### Queue

for FIFO operations

### Maps

idr for ID allocation

↳ Red-black tree

↳ self balancing BT for efficient  
lookups → (CFS)

## Interrupts and Interrupt handler

Top half → immediate  
bottom half → minimal work.

## Interrupts

↳ run in atomic context, can not sleep  
or access user-space.

## Kernel Synchronization

→ Spinlocks → for shorter critical section.

→ Mutex → for longer sections.

→ Semaphore.

↳ Counting mech for resource access.

## Timer and Timer Management

### Spinlocks

## Memory Management

→ Pages → physical memory divided  
4 KB

zones

↳ DMA, NORMAL, HIGHMEM.

→ Page allocator.

↳ large chunks.

→ Slab allocator

↳ Cache for frequently used  
objects.

## Virtual memory

address space managed via struct  
mm\_struct.

→ Technical Deep Dive

→ Kernel

\* Interface b/w user programs and hardware.

It is also the core part of the OS.

In Linux the Linux OS is usually referred to as kernel.

→ Subsystems

Process Management

- handles process lifecycle and scheduling of all processes.

Memory Management Unit

- handles memory allocation, physical memory or virtual memory.

Virtual File System

- Provides more like API or syscall for physical storage device.

Networking Stack

- responsible for how a user processes data packets from user space all the way down to the kernel.

\* Ubuntu

↳ upstream kernel with modifications

everyone

contribution

Device Driver

- handles different devices in the kernel → or attached to the system

Arch Specific logic

- processor specific logic present in the kernel tree.

Modifying the Linux kernel

(C, Rust)

enable/disable features

optimise kernel performance and size

add support for new hardware.

apply security features.

\* make menuconfig (GUI) ...

\* make defconfig (default configuration)

- 1) Start by building your kernel. (make -jx)

- 2) Install the kernel.

- 3) Running your kernel.

Test and debug

\* What is the change you made.

↳ dmesg log

↳ Captures everything coming out of the kernel.

## Automated testing

→ **linux test Project (LTP)** ↳ 3-4 hrs.

→ **kselftest (test)** ↳ part of linux kernel

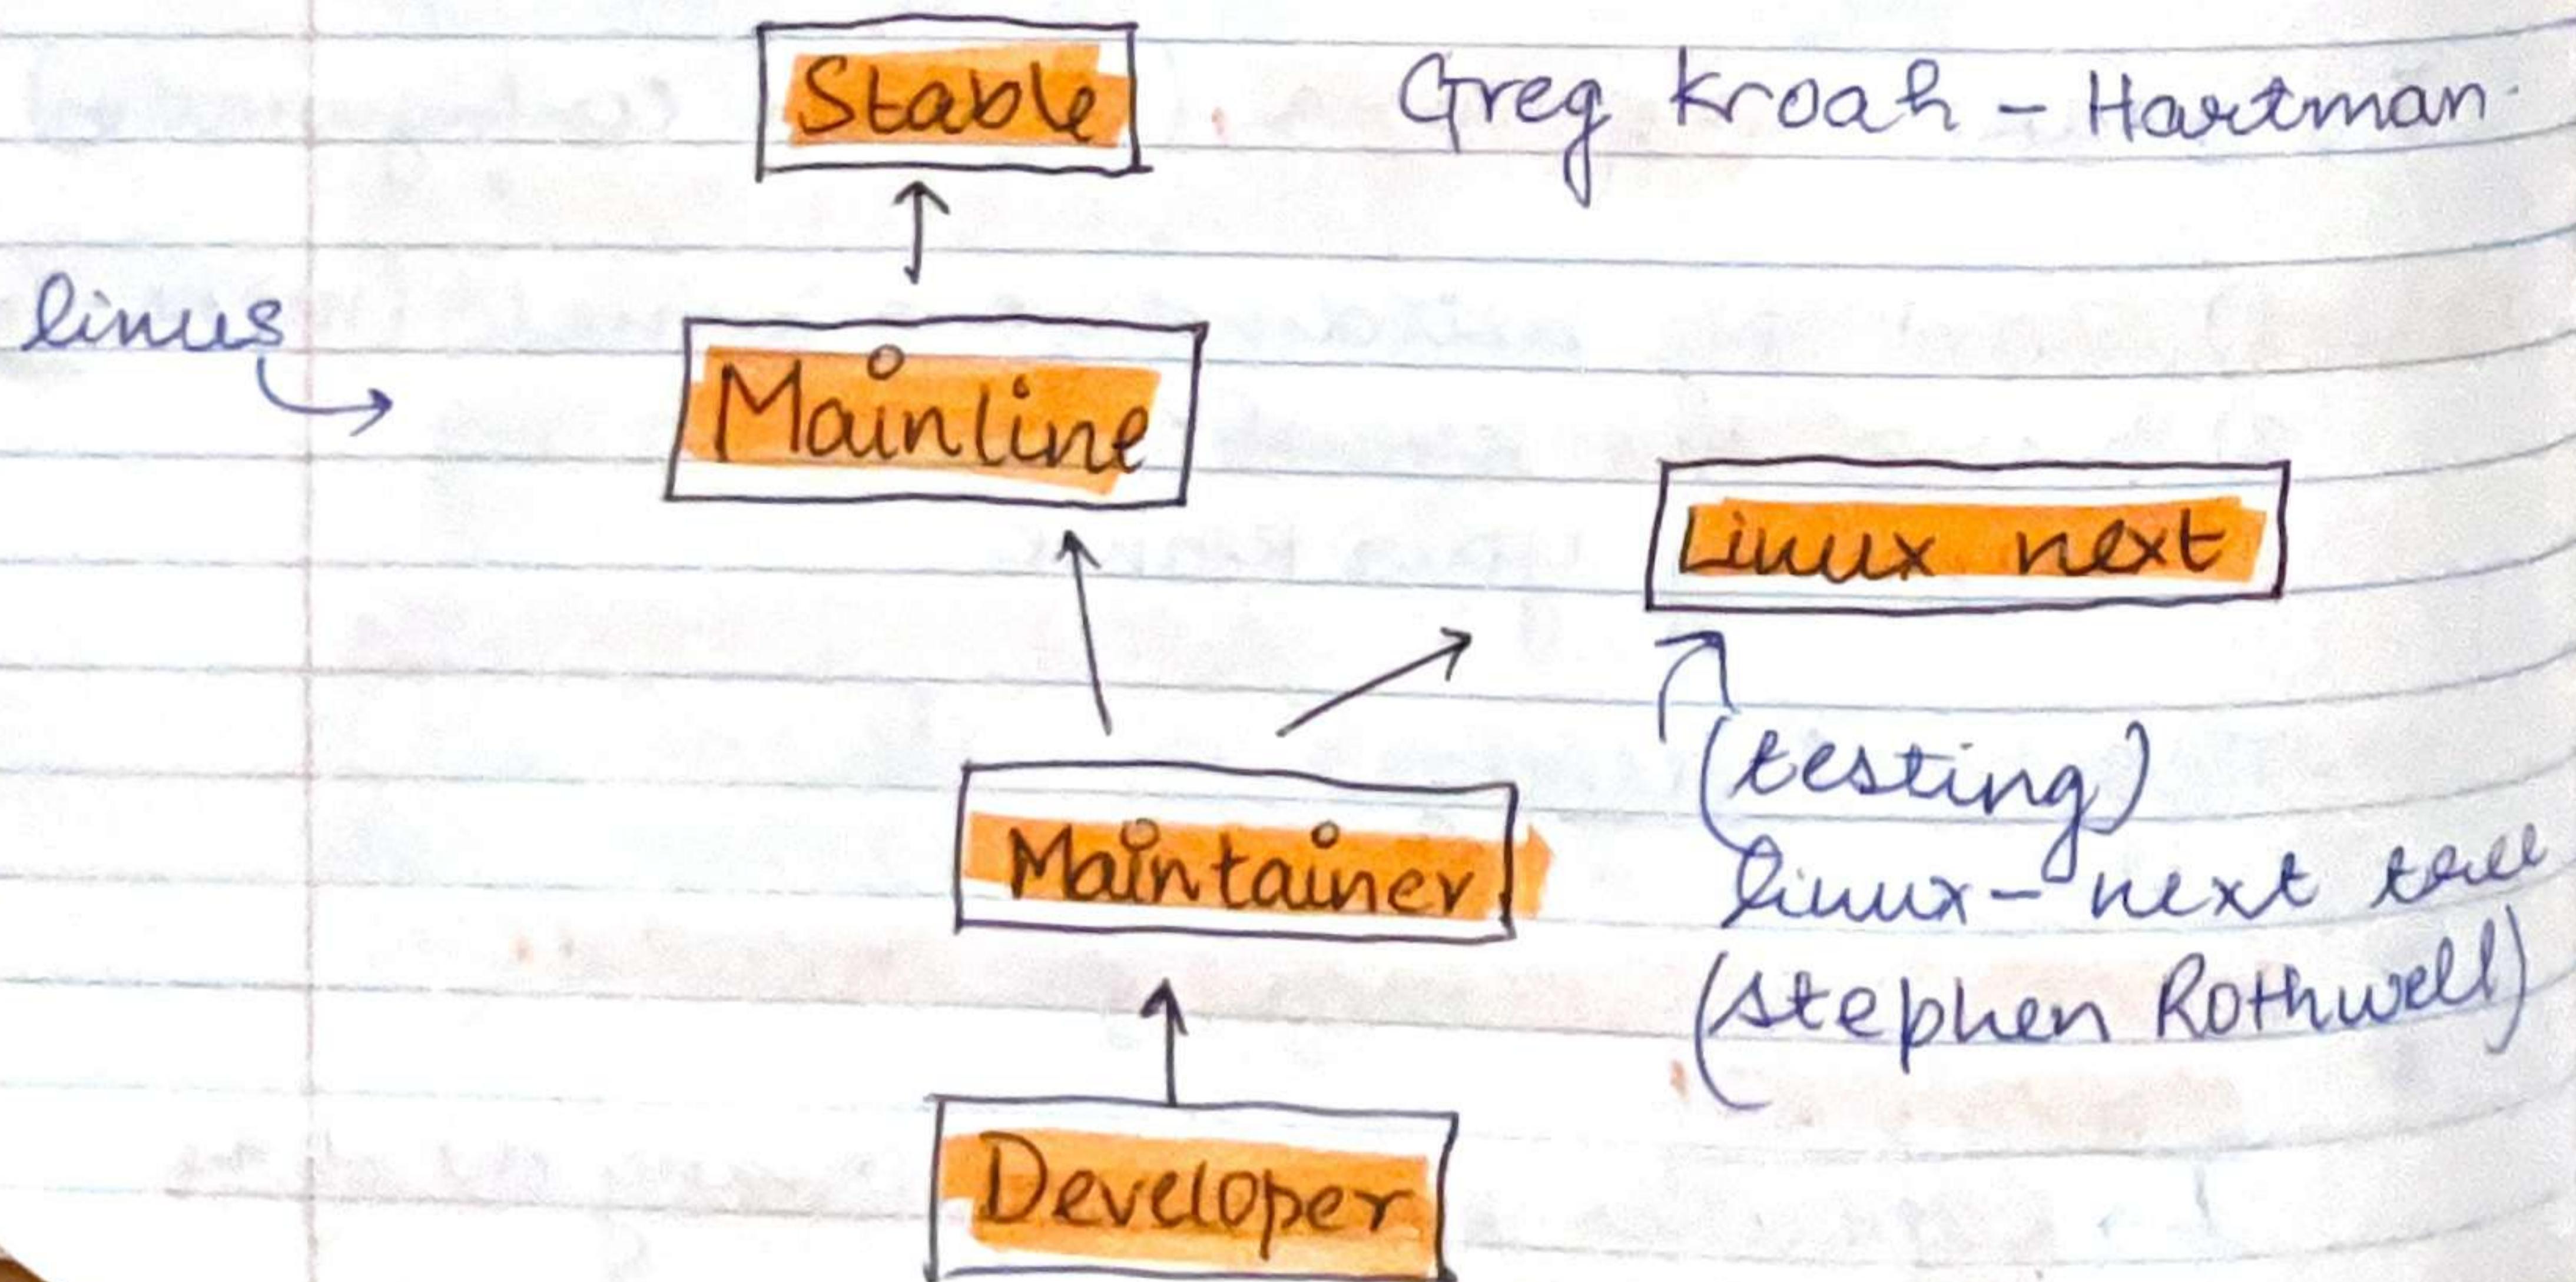
Imp

debugging · **printk()**

• **trace** : trace of which func called this one

• **dump stack**: prints stack tree

## linux kernel development Process



\* The linux kernel is both monolithic and modular since it can insert and remove kernel modules at runtime

### \* Monolithic Kernel

↳ All the kernel code is running in a single address space (or) kernel code running in kernel space

### Microkernel

↳ Main kernel code is running in kernel space and DD and modules are running in user space

Only load (files) when needed

### Hybrid

linux (current) → Monolithic  
+  
you can add DD & modules

menuconfig

→ kernel config

M3 → loadable kernel module

→ runtime.

## Steps to configure and compile linux kernel

- 1) get linux source code (git)
- 2) install dependencies. (apt update)
- 3) choose a default configuration. (defconfig)
- 4) clean files of old dependencies (make clean)
- 5) configure the kernel (make menuconfig)
- 6) compile the kernel (make -j\$(nproc))
- 7) make modules
- 8) Install the kernel (sudo make modules\_install)
- 9) ~~Reboot~~ - update GRUB.
- 10) Reboot.

Q. What is the role of OOM killer.

ans: Out - of - memory killer is a feature in linux kernel designed to handle situation when the system runs out of available memory.

Terminates process by weighing pro's and con's

→ less critical

→ uses most memory.

X Q. What is kernel panic?

ans: • fatal error that linux kernel encounters and recover from.  
• Happens when the kernel detects an inconsistency or an issue that severely impact system stability.  
• Often requires manual restart.

Q. What is loadable kernel module (LKM)?

ans: Specific component of kernel that can be dynamically loaded and unloaded runtime.

Doesn't require system to be rebooted.  
to add or remove at runtime.

Q. How does the kernel handle security and permissions?

ans: Uses mechanisms like user account, access control list and security module.

→ Makes sure that users and processes access resources they are authorized to use.

Q. What is Kernel Preemption.

ans: Ability of kernel to interrupt a running process to switch to a higher priority task.

Imp because gives kernel access to handle critical tasks first.

Q. How does kernel enforce memory protection between different processes?

ans:

- (1) Virtual memory and page table.
- (2) Privilege levels
- (3) COPY - on - write.

Optimise memory sharing b/w forked processes.

(4) Page Permissions

(5) Segmentation faults

↳ block unauthorized access to memory.

(6) IPC

✓/5) Role of kernel managing i/p and o/p operations.

ans:

~~Disk scheduling algorithms~~

↳ FCFS

↳ SSTF

↳ SCAN end

↳ LOOK last read

↳ C-SCAN

↳ C-LOOK

ans: By using system calls, device drivers to access HW. VFS, disk scheduling, buffering and caching, interrupt handling and direct memory access.

Q Handling System Calls in multithreaded environment?

ans: if a system call blocks, only the calling thread is blocked, not the entire process.  
Some system calls are thread safe that means they can be executed by multiple threads.

eg: `Read()`

Imp

Q How does the kernel handle real-time processing requirements.

ans: By using RT scheduling policies  
Rate monotonic scheduling and Earliest deadline first.

Q How does the kernel ensure data integrity in file systems.

ans: kernel uses mechanism such as journaling, checksum, redundancy.

RAID  
type stuff.

Imp

Q How does kernel secure a system against Vulnerabilities?

ans: (1) access control  
(2) enforcing memory protection.  
(3) isolating resources.  
(4) SE Linux  
(5) AppArmor.

= Kernel updates and patches.

= Kernel address space layout randomization

Q How does kernel handle hardware interrupts. reduce int latency

ans: Interrupt handler are used, which are special function that execute in response to an interrupt signal. They ensure timely processing of events.

Interrupt Descriptor Table and how the kernel maps interrupts to their resp handler.

## \* Kernel → Kernel Arch.

mb

Q. How does kernel manage process isolation and resource allocation in containers?

ans: namespaces and cgroups

↳ isolation

↳ of

↳ processid

network interface

file system.

cgroups

↳ control resource utilisation

Q. How does kernel handle priority inversion.

ans: Priority Inheritance

↳ low priority task gets high priority

Priority ceiling

↳ resource access restricted to higher priority task.

Q. Fork() bomb?

ans: Fork bomb ~~sometimes~~ spawns excessive child processes, depleting system resources.

Mitigate this by ulimit -u

Q. Device hotplugging?

ans: The kernel allows

This allows hardware component to be added or removed while the system is running without requiring boot.

Kernel handles this using

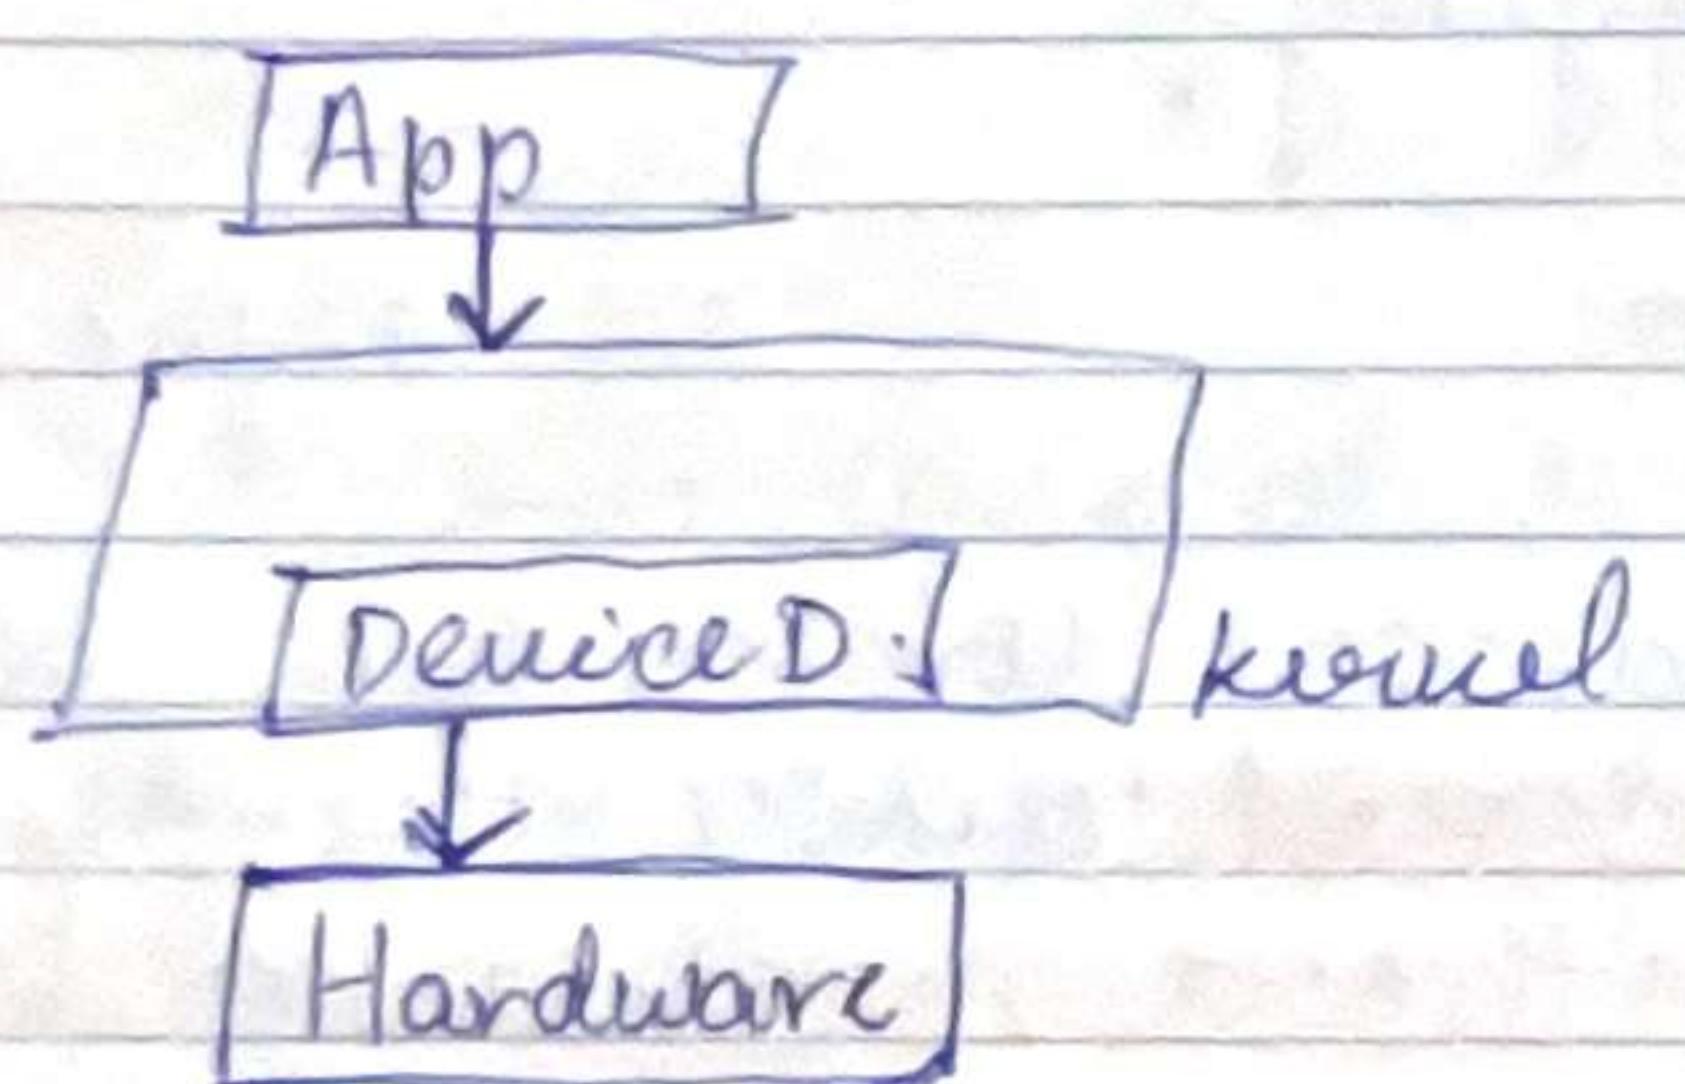
dynamic device detection  
device loading  
resource allocation

## → Completely fair Scheduler

- \* It models an ideal, precise multitasking CPU.

## Device Drivers

- \* They are abstraction to a piece of hardware.



## Types of Drivers

- 1) Character ✓
- 2) Block ✓
- 3) Network ✓
- 4) Pseudo ✓
- 5) USB ✓
- 6) Platform ✓
- 7) File Systems ✓

### 1) Character Device Driver

used for: Device that handle data as a stream of bytes.

eg: Serial ports, keyboards

Interact with kernel:

- register\_chrdev() func to register the driver
- file operations  
read(), write(), open(), close()

eg: /dev/ttys0  
↳ char devices.

### 2) Block Device Driver

used for: Devices that store and transfer data in fixed size block

eg: SSD, USB device

Interaction with kernel:

- register\_blkdev() func to register the driver
- used block request queue to manage read and write op.
- Works with VFS and I/O Scheduler

eg: /dev/mmcblk0 (SD card)

### 3) Network Devices.

used for: Network interface like ethernet, wifi, bluetooth adp.

#### Interaction

- register with register\_netdev()
- Struct net\_device interface
- socket buffer for packet transmission

eg: eth0.

### 4) Pseudo

used for: virtual devices that don't correspond to real hardware

#### Interaction

- implements file operation read(), write()
- often used for testing and debugging

### 5) USB

used for: USB peripherals like flash drives, printers, webcam.

#### Interaction

- uses usb\_register\_driver()
- communication using ~~URB~~ USB Request Block.

### 6) Platform.

used for: on-board SOC peripherals, GPIO controller, SPI/I2C dev

#### Interaction

- Device tree (.dts)
- registered with platform bus via platform\_driver\_register()

### 7) File System.

used for: Managing file systems like ext4, NTFS.

#### Interaction

- Registers with the VFS via register\_filesystem()
- implement inode op (read\_inode(), write\_inode())

eg: proc, sys, mnt

IMB

How devices interact with the kernel.

- 1) Drivers must register with the kernel using API.

register\_chrdev()

register\_blkdev()

usb\_register\_driver()

- 2) Handling System Calls

- 3) Managing hardware via kernel subsystems

I/O Scheduler → block devices

network stack → network devices

USB core → USB devices

Device tree → platform devices.

- 4) Interrupt handling

Devices generate interrupts and the driver must register an interrupt handler to handle them asynchronously

- 5) Power Management

Embedded drivers need to support suspend/resume operation for power efficiency

Building device driver / Modules

→ Hello World

printf() → used to print in kernel  
seen in logs

Syslog

↳ System log has the print for  
Hello World

'tail' / var / log / syslog

ls - l

Show which device file is a  
char device or a block device

This command also gives you a  
major or minor number

Major, Minor.

Drivers

1, 4, 7, 10

The devices

refer to the  
device

/dev/null/driver1

/dev/zero

serial 3 4

VCS1 & VCSA 7

Every Module has some operation written in it.

Struct file-operation Scull-fops {

- Owner
- lseek
- read

:

}

each field, points to the function in the driver that implements a specific operation

Implements function for typical system call

write()

↓  
doesn't immediately write to it.

It is stored in the buffer and it is written once file is being closed or buffer is being flushed.

Q Explain the role of ioctl() in linux device driver?

ans:

ioctl() is a system call ~~takes~~ to provide generic interface for controlling and configuring a device.

It is used to implement device specific functions that are not part of standard system calls.

Takes 3 arguments.

★ Q

How do you implement a character device driver in linux?

ans:

- a) define the device structure (operations, name)
- b) implement the device file operations
- c) Register the device
- d) Create the device file
- e) Implement device specific func.
- f) Test the driver

★ Q

Describe the process of allocating memory in linux device driver?

ans: kmalloc() → allocate memory from kernel heap

vmalloc() → allocate in virtual memory space

X

Q. Describe the process of cross compiling a device driver for embedded linux?

Ans: Install cross compiler.

↳ used to build code for target architecture.

↳ Configure the build env.

↳ compatible with cross compiler and target architecture.

↳ Build driver

↳ build and update and compile.

↳ Test driver

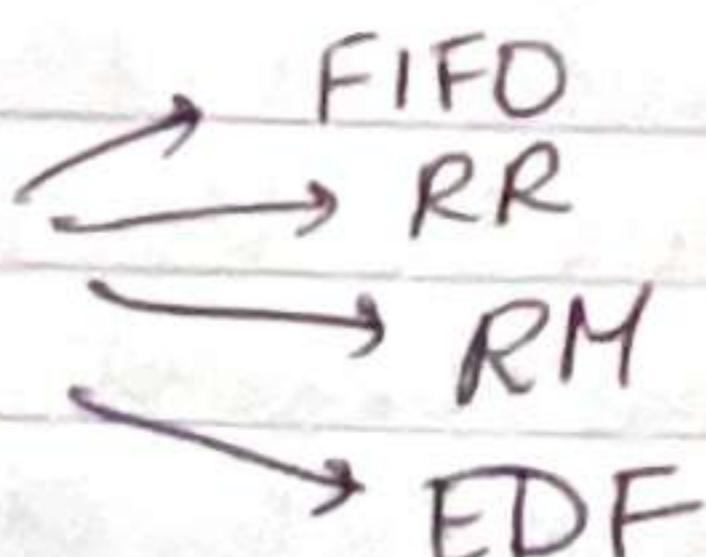
↳ on target system.

Q. How do you handle real-time constraint in an embedded linux device driver.

Ans: (a) interrupts

(b) DMA

(c) Task scheduling



d) locking → for process synchronisation.

Q. SPI

12C

\* 4 wire

~~SCK~~, SCK

MISO, MOSI

\* 2 wire

SDA, SCL

\* full duplex

\* fast

\* displays,

memory, wireless  
transceiver

\* half duplex

\* slow.

\* sensors, EEPROM

real-time clock.

Q. Describe the use of GPIO in embedded linux device driver.

Ans: GPIO device driver uses Linux kernel's GPIO framework, which provides a consistent interface for device driver to interact with GPIO pins.

- Allocate/deallocate GPIO pins
- Setting the pin direction, value
- registering in interrupt handler

## Snap vs AppImage vs Flatpacks.

Q Optimise performance of system call?

- ans:
- (a) Minimise system calls.
  - (b) use interrupt handler
  - (c) use DMA.
  - (d) use better data structures.
  - (e) Optimise memory.

### Package

→ dpkg

low level package manager.

(cons.)

① doesn't install dependencies.

- download the .deb file

- sudo -i discord.deb

↳ fails

because of dependencies.

dpkg -l : list all packages.

~~highlevel~~

→ apt (Advanced Package Tool)

sudo apt install <name-pkg>

if anything is broken (dependencies)

sudo apt --fix broken install.

### APT

has a repository

↳ collection of software.

list

Show <pkg>

remove

→ remove app but not config

list -- installed

purge → remove everything

update → up-to-date list/Repo

upgrade → upgrade the list

full-upgrade → remove previously installed & not upgraded

high-level

→ Aptitude (GUI)

More like apt on steroids

→ Snap

Snap → uses store.

makes it easier to get apps on-to a repository. store.

getting it onto a repo is difficult.

→ language based / git

Pip

Git

git clone < Repo url >

Pip

sudo apt install python3-pip

pip3 -r req

SNAPS

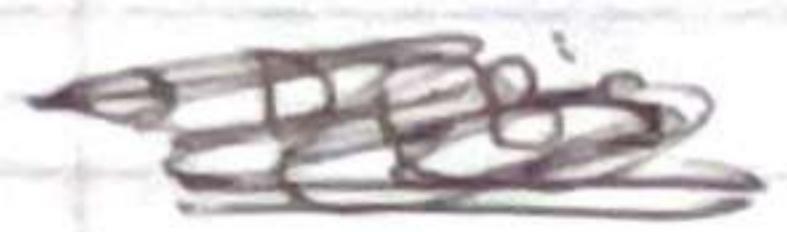
Build once

distribute everywhere

↳ do not have to make  
PKG for every distro  
or arch

— Auto update

— Containerized!!!



## → Snap sandbox

Allows sharing of data & unix  
sockets between snaps.

Why?

Share common libraries  
reduce size by avoiding duplication.

uses AppArmor by default.

↓  
(LSM)

linux security module

Other distro like Fedora

✓  
SELinux

→ degrade SNAPS. Sandbox.

## Security

### → Linux Security Module

They provide MAC, mandatory  
access control.

MAC : limits access based on what  
system user wants or defines

DAC : Security policies that gives user  
chmod access to decide who to grant  
permissions to.

file permission eg: google docs

1) SE Linux MAC, RBAC.  
Security enhanced linux

- ↳ developed by NSA
- ↳ default in Redhat, Fedora

uses labels to define policies for every  
file, process and user

3 modes

- ↳ enforcing → block action
- ↳ Permissive → log but allow action
- ↳ disabled

\* complex

\* can cause unexpected ~~denial~~ denial

Deny everything

→ embedded → lightweight

## 2) AppArmor MAC, RBAC,

↳ Developed by Canonical (Ubuntu)  
↳ debian and ubuntu.

uses profiles to define rules for application.

Profiles stored in /etc/apparmor.

- \* easier
- \* less fine grain control

## Q. What are LSM; how they work?

ans: LSM are security module that enforce MAC ~~at the R~~

LSM hook into linux kernel and enforce security policies.

allows everything

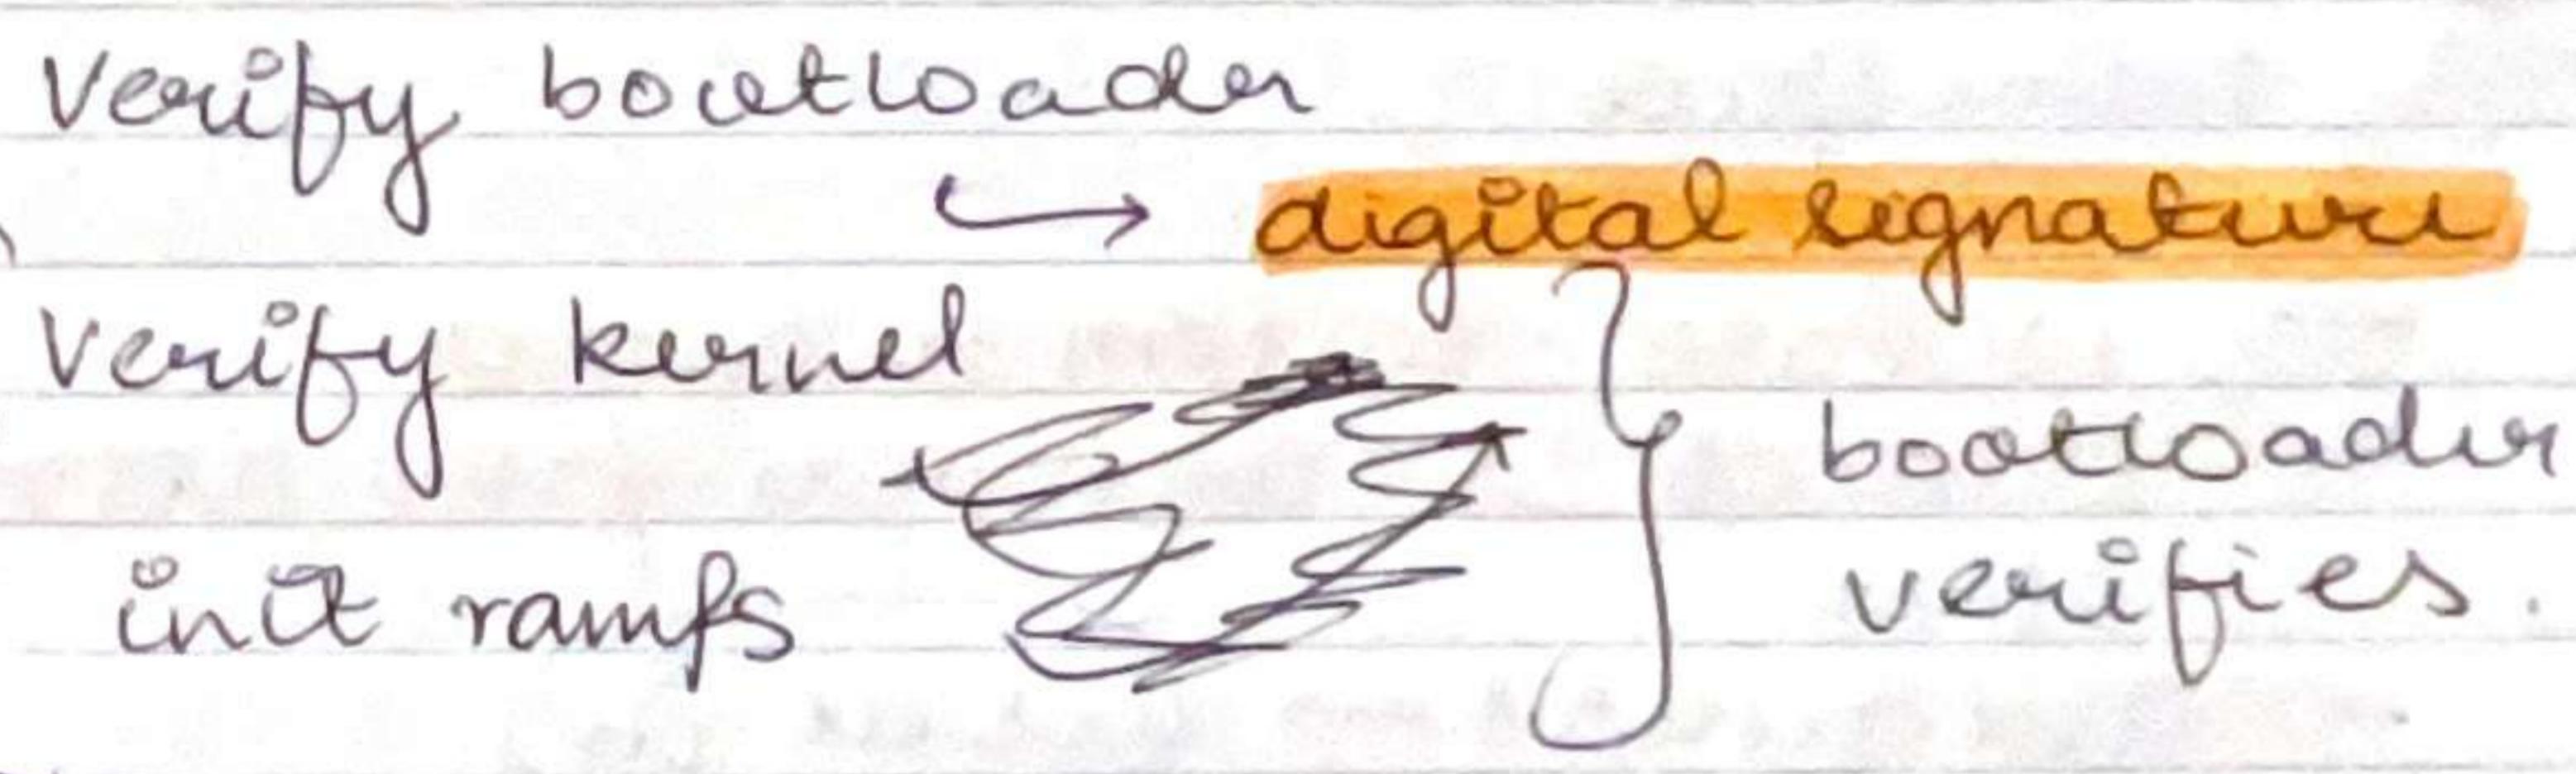
## Q. Securing an embedded linux system.

ans:

- remove unwanted access
- LSM → to implement MAC
- no root ssh access
- asymmetric encryption
- disk encryption
- secure boot

## Q. What is secure boot?

ans: Secure boot is a security feature that ensure cryptogra... signed bootloader and kernels are executed



prevents kernel modification,  
bootloader tampering

Q. Security risks of loading kernel module?

ans: unsigned module → malicious code  
kernel crash → kernel code not written well.

Backdoor into kernel → kernel modules can even undetected

\* (restrict who and what module can be loaded).

Q. linux namespace and cgroups.

namespace

- isolate system resource.
- prevents 1 container from seeing the other.
- containers → used in

eg: Process in 1 namespace can't see the other.

Cgroups

- control resource usage.
- prevents resource exhaustion attacks.
- used in container, systemd services

eg: limit memory usage of 1 container.

Q. What is stack Canary Protection?

How does linux prevent buffer overflow

Ans: Stack Canaries are random values placed on stack to detect buffer overflows before they can overwrite return address.

- \* random value placed b/w local and return add.
- \* before returning from a function the canary value is checked.
- \* if value is changed due to overflow then abort execution.